

Solution 2

Initial solution (TOO SLOW)

What I wrote first	Why it hurt
<code>unordered_set<pair<int,int>> seen</code> – test “visited?”	Hashing every coordinate ($\approx 1\,000\,000$ times) + pointer chasing.
<code>unordered_map<pair<int,int>, int> coord2id</code> – map cell \rightarrow island-id	Same cost again for every insertion / lookup.
<code>unordered_map<int,int> id2paint</code> – island-id \rightarrow picture count	Extra hash layer for something that is really an array.
<code>string</code> grid \rightarrow <code>museum[i][j] = s[j]</code>	Millions of bounds checks on short strings.

On the maximal 1000×1000 board the BFS visits $\approx 10^6$ cells.

Each visit did **three** hash operations and several heap indirections – more than 30 million hash probes total.

Run-time $\rightarrow \sim 1.3$ s (local) \rightarrow **TLE** on the judge.

Better solution (ACCEPTED)

1. One dense 2-D array instead of two hash tables

```
vector<vector<int>> id(n, vector<int>(m, -1)); // -1 = not visited
```

- `id[x][y] == -1` \rightarrow “not seen yet” (so `seen` set deleted).
- The value itself is the island number (so `coord2id` deleted).

1. Plain vector for island data

```
vector<long long> paint; // push_back once per island
```

- Flood-fill** each island once, store its picture count in `paint[id]`.
- Answer queries** in $O(1)$:

```
cout << paint[ id[x-1][y-1] ] << '\n';
```

1. Moved the `dx[]` , `dy[]` arrays outside the loop to avoid rebuilding them.

Why it worked

metric	before	after
hash look-ups per cell	3-4	0
peak heap	≈ 30 MB	≈ 6 MB
run-time 1000×1000	> 1 s (TLE)	≈ 0.02 s (AC)

Array indexing is two orders of magnitude cheaper than hashing and has perfect cache locality. Once the big hash tables were gone, the exact same algorithm (BFS, $O(n \cdot m + k)$) flew under the time limit with plenty of margin.